


# CMPT 732-G200. Practices for Visual Computing

Ali Mahdavi Amiri

# AI, ML, DL

- Artificial Intelligence (AI)
- Machine Learning (ML)
- Deep Learning (DL)




# AI

- AI was born in 1950s.
- Can computers think?




# AI

- The effort to automate intellectual tasks normally performed by humans.



# Symbolic AI

- For a long time, experts believed that human-level artificial intelligence could be achieved by having programmers **handcrafted** a sufficiently large set of **explicit rules** for manipulating **knowledge**.





# Symbolic AI

- For a long time, experts believed that human-level artificial intelligence could be achieved by having programmers handcrafted a sufficiently large set of **explicit rules** for manipulating **knowledge**.
- It gets intractable in many tasks (e.g., image classification)



# Machine Learning

- Lady Ada Lovelace with Charles Babbage introduced Analytical Engine which was a mechanical computer.



# Machine Learning

- As mentioned by Ada Lovelace, this machine was limited to the tasks that were instructed by human (programmers).






# Machine Learning

- Can machines go beyond instructions?
- Can they become **creative**?





# Machine Learning

- These discussions lead to introducing Turing Test in 1950.





# Machine Learning

- As opposed to classical programming, in machine learning, a set of rules are **learned**.





# Machine Learning

- As opposed to classical programming, in machine learning, a set of rules are **learned**.
- These rules can be applied to **new data** to produce **original** answers.



# Machine Learning


- As opposed to classical programming, in machine learning, a set of rules are **learned**.
- These rules can be applied to **new data** to produce **original** answers.
- In fact, machines are **trained** rather than directly programmed.





# What is Learning?

- Components of Machine Learning:
  - **Input** data points
    - An image of a digit



# What is Learning?

- Components of Machine Learning:

- **Input** data points

- An image of a digit



- Examples of the **expected output** (Ground Truth)

- A set of digits and their labels **5,0,4,1**

# What is Learning?

- Components of Machine Learning:

- **Input** data points

- An image of a digit



- Examples of the **expected output** (Ground Truth)

- A set of digits and their labels **5,0,4,1**

- Measurement of the **performance** of the machine


- If machine returns the correct label for a digit.

# What is Learning?

- Components of Machine Learning:

- **Input** data points

- An image of a digit



- Examples of the **expected output** (Ground Truth)

- A set of digits and their labels **5,0,4,1**

- Measurement of the **performance** of the machine
    - If machine returns the correct label for a digit.

# What is Learning?

- Measurement of the **performance** of the machine
  - If machine return the **correct label for a digit**.
- This process gives a feedback to **adjust** the way algorithm works.



# What is Learning?

- Measurement of the **performance** of the machine
  - If machine return the **correct label for a digit**.
- This process gives a feedback to **adjust** the way algorithm works.
- This adjustment is **learning**.

# Models

- Machine learning models can learn in two main forms
  - Unsupervised
  - Supervised

# Unsupervised Learning

- Learning a function that learns a commonality in data that has not been classified/labeled/categorized.

# Unsupervised Learning

- Learning a function that learns a commonality in data that has not been classified/labeled/categorized.
- We do not have ground truth data.

# Unsupervised Learning

- Learning a function that learns a commonality in data that has not been classified/labeled/categorized.
- We do not have ground truth data.
  - K-means
  - PCA
  - SVD




# Unsupervised Learning

- Learning a function that learns a commonality in data that has not been classified/labeled/categorized.
- We do not have ground truth data.
  - K-means
  - PCA
  - SVD


# K-means

- Unsupervised model to cluster  $n$  data point




# K-means

- Unsupervised model to cluster  $n$  data point into  $k$  clusters.



# K-means

- Main objective is to minimize




$$\sum_{i=1}^k \sum_{x \in S_i} \|x - \mu_i\|^2$$

$i$ th cluster

# K-means

- Main objective is to minimize



$$\sum_{i=1}^k \sum_{x \in S_i} \|x - \mu_i\|^2$$

$\swarrow$   
 $i$ th cluster

$\searrow$   
Mean of points in  $S_i$



# Algorithm


- Assignment step: Assign each point  $x$  to the cluster  $S_i$  whose mean  $\mu_i$  is closest to  $x$ .

# Algorithm

- Assignment step: Assign each point  $x$  to the cluster  $S_i$  whose mean  $\mu_i$  is closest to  $x$ .
- Update step: Calculate new means to the centroids of data points in new  $S_i$ .


# Algorithm

Assign initial centroids




# Algorithm

Assign initial centroids




Assignment step




# Algorithm


Assign initial centroids



Assignment step




Update step




# Algorithm


Assign initial centroids




Assignment step



Update step



Assignment Step



# Supervised Learning

- Learning a function that maps an input to an output based on example input-output pairs.

# Supervised Learning

- Learning a function that maps an input to an output based on example input-output pairs.
- We have ground-truth data.



# Supervised Learning


- Learning a function that maps an input to an output based on example input-output pairs.
- We have ground-truth data.
  - Least square line
  - Linear regression
  - Deep Neural Network (supervised)

# Supervised Learning

- Learning a function that maps an input to an output based on example input-output pairs.
- We have ground-truth data.
  - Least square line
  - Linear regression
  - Deep Neural Network (supervised)


# Least Square Lines

- Given data  $\{(x_1, y_1), \dots, (x_N, y_N)\}$ , we want to find a line  $y = mx + b$  which is close to our data.



# Least Square Lines


- Given data  $\{(x_1, y_1), \dots, (x_N, y_N)\}$ , we want to find a line  $y = mx + b$  which is **close** to our data.



$$E(a, b) = \sum_{i=1}^N (y_i - (mx_i + b))^2$$

# Least Square Lines


- Given data  $\{(x_1, y_1), \dots, (x_N, y_N)\}$ , we want to find a line  $y = mx + b$  which is **close** to our data.



$$E(a, b) = \sum_{i=1}^N (y_i - (mx_i + b))^2$$

# Least Square Lines

- Given data  $\{(x_1, y_1), \dots, (x_N, y_N)\}$ , we want to find a line  $y = mx + b$  which is **close** to our data.




$$E(a, b) = \sum_{i=1}^N (y_i - (mx_i + b))^2$$

Minimize this squared distance


# Least Square Lines

- What should we do?



# Necessities


- Data






# Necessities

- Model



# Necessities

- Loss function (differentiable)




$$E(a, b) = \sum_{i=1}^N (y_i - (mx_i + b))^2$$

# Neural Networks

- Single lines are too restrictive. We want a more general non-linear model.


# Neural Networks

- Trainable functions that can be used as a mapping between an input and a desired output.




# Neural Networks

- Trainable functions that can be used as a mapping between an input and a desired output.




# Neural Networks

- Trainable functions that can be used as a mapping between an input and a desired output.




# Neural Networks

- Neurons are connected to each other to define a parametric function.



# Neural Networks

- The value input to each neuron is found by a linear combination of neuron values.




$$z = w_0a + w_1b + w_2c$$



# Neural Networks


- You can also add a bias to each node



$$z = w_0a + w_1b + w_2c + b_0$$

# Neural Networks


- Followed by a non-linear activation function.



$$z = w_0a + w_1b + w_2c$$

# Activation Function


- Sigmoid



$$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}}$$

# Activation Function


- Hyperbolic Tangent



$$f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

# Activation Function


- Rectified Linear Unit (ReLU)



$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$

# Activation Function

- Leaky Rectified Linear Unit (ReLU)



$$f(x) = \begin{cases} x & \text{for } x > 0 \\ \alpha x & \text{else} \end{cases}$$

# Training

$$\textit{Min} \sum_{i=1}^n \mathcal{L}(f(x_i, W), y_i)$$

# Training

$$\text{Min} \sum_{i=1}^n \mathcal{L}(f(x_i, W), y_i)$$

Loss function





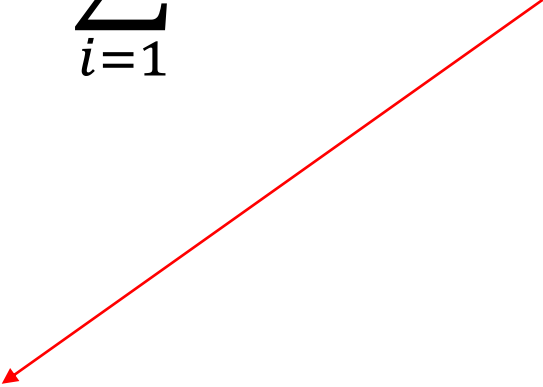
# Training

$$\text{Min} \sum_{i=1}^n \mathcal{L}(f(x_i, W), y_i)$$

Output of neural network

# Training

$$\text{Min} \sum_{i=1}^n \mathcal{L}(f(x_i, W), y_i)$$



Expected result or ground truth

# Training

- We use gradient descent to find the minimum

$$\textit{Min} \sum_{i=1}^n \mathcal{L}(f(x_i, W), y_i)$$

# Training


- We use gradient descent to find the minimum

$$\text{Min} \sum_{i=1}^n \mathcal{L}(f(x_i, W), y_i)$$

- We update the weights little by little based on partial derivatives

$$W_{t+1} = W_t - \alpha \nabla_w \mathcal{L}$$

Learning Rate



# Training

- We use gradient descent to find the minimum

$$\text{Min} \sum_{i=1}^n \mathcal{L}(f(x_i, W), y_i)$$

- We update the weights little by little based on partial derivatives

$$W_{t+1} = W_t - \alpha \nabla_w \mathcal{L}$$



Defining a meaningful loss function is important

# Training

- We have a set of images and we want to train a model on these images so that the model **generalizes** well.

# Training

- We have a set of images and we want to train a model on these images so that the model **generalizes** well.
- It performs well on data sets that are not in the training set.

# Training

- Having a data set, we split it into three groups:
  - Training set
  - Test set
  - Validation set

Why Three Sets? Why not only Training set and Test set?



# Training

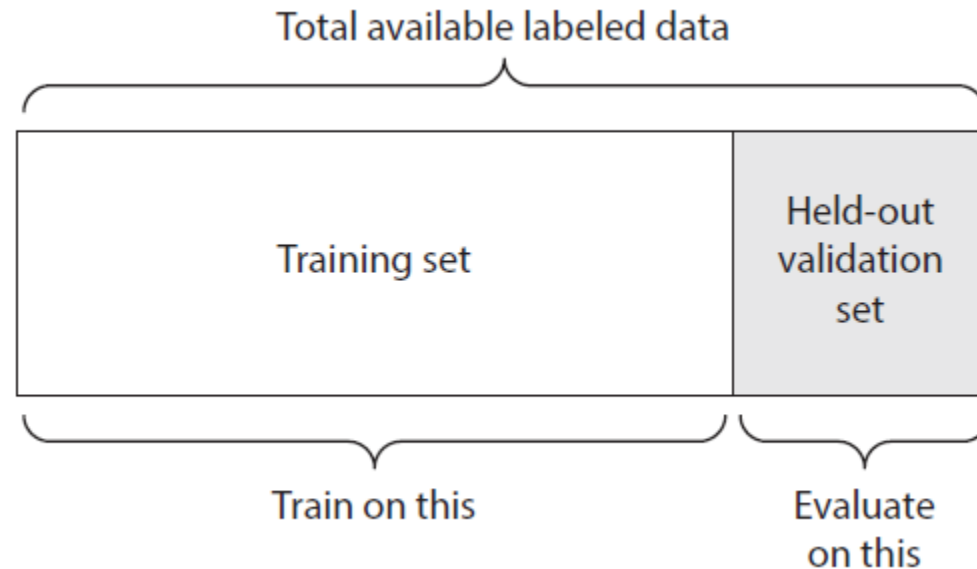
- You train your model using training set and **fine tune** the model (hyperparameters) using **validation set**. Each time you fine tune, a little bit of information **leaks** to the network. Although, you have not trained your model on validation set, the model overfits to the validation set.

# Training

- You train your model using training set and fine tune the model (hyperparameters) using validation set. Each time you fine tune, a little bit of information leaks to the network. Although, you have not trained your model on validation set, the model overfits to the validation set.
- You need an extra set to actually measure the performance of your network that has been **never exposed** to the network **directly** or **indirectly**.


# Validation

- **Hold-out validation**: you put aside a portion of your training data for validation.



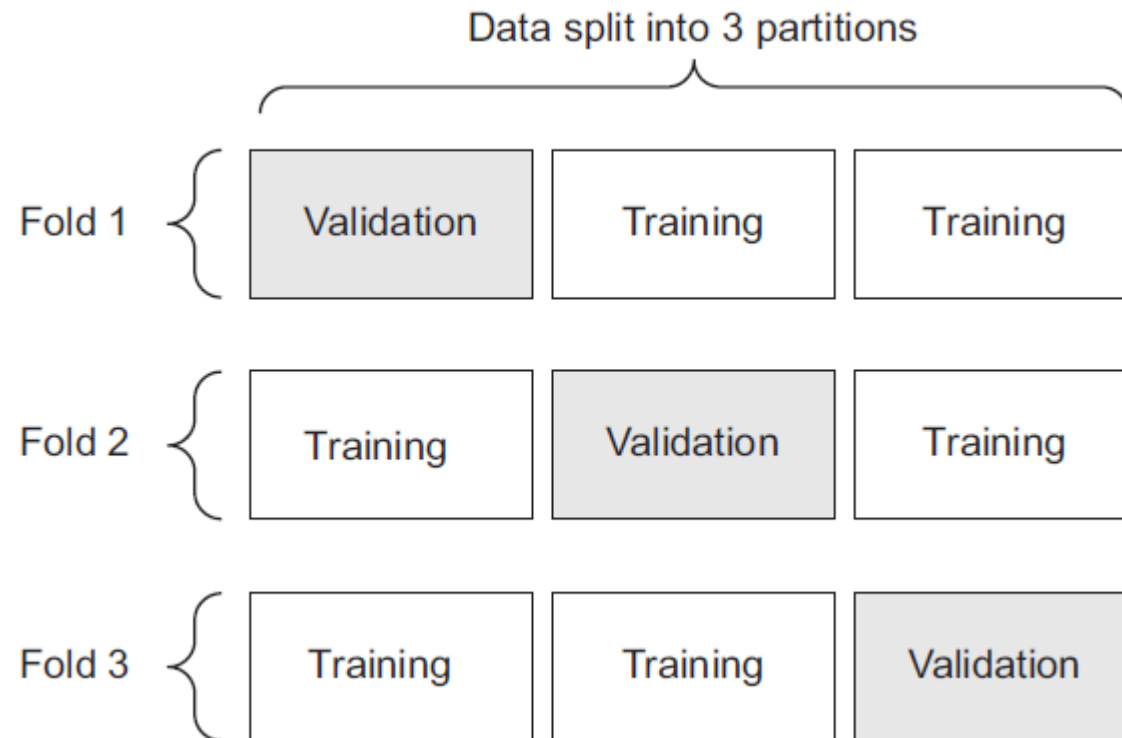
# Validation

- **K-Fold Validation:** Split the data into k Partitions of equal size.



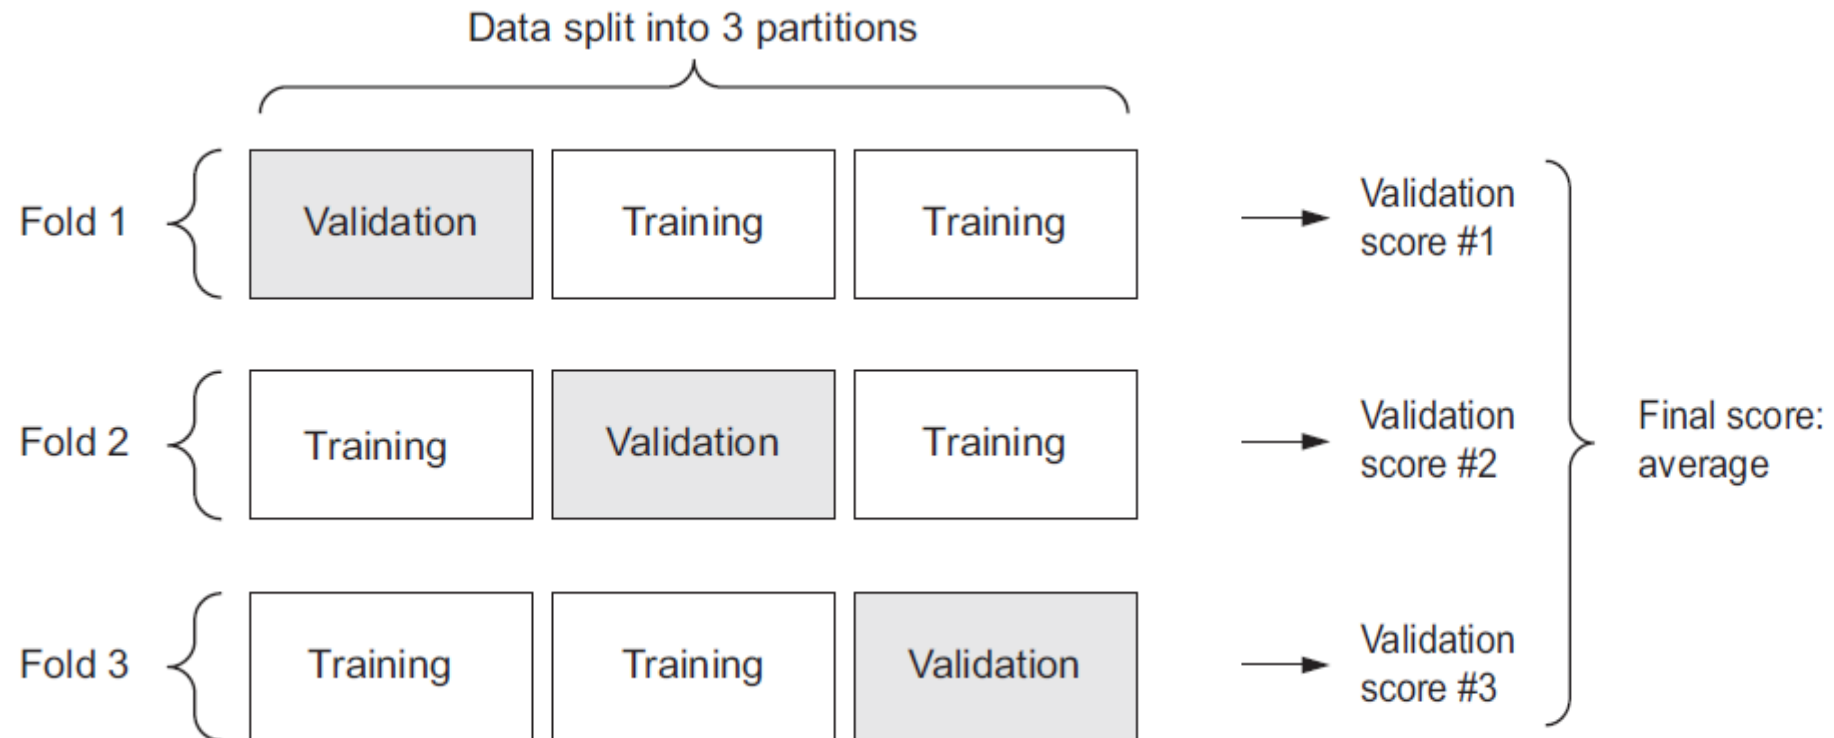
# Validation

- **K-Fold Validation:** Validate on  $i$ th partition, train on the rest.




# Validation

- **K-Fold Validation:** Final score is the average of all the scores.



# Data Processing

- **Vectorization**: Neural networks usually accept vectors, you will always need to vectorize your input data.



# Data Processing

- **Data Normalization**: You should avoid providing large data values or heterogeneous data values to your network.



# Data Processing

- **Data Normalization**: You should avoid providing large data values or heterogeneous data values to your network.
- **Large Values**: You will end up having **large gradient updates** and the network does **not converge**.

# Data Processing

- **Data Normalization**: You should avoid providing large data values or heterogeneous data values to your network.
- **Large Values**: You will end up having large gradient updates and the network does not converge.
- **Heterogeneous Data**: Some data have more effect on the network than the others.

# Data Processing

- Rule of thumb:
  - Normalize each feature independently to have mean 0.
  - Normalize each feature independently to have standard deviation of 1.

# Handling Missing Values

- If you have missing values in some features, default them to be zero. Network will learn to ignore such values.

# Feature Engineering

- Try to provide data sets that are easier to understand by the network.

# Feature Engineering

- Pixel values of a clock are a lot harder to understand for a network rather than a simple numeric representation.

Raw data:  
pixel grid



---

Better  
features:  
clock hands'  
coordinates

{x1: 0.7,  
y1: 0.7}  
{x2: 0.5,  
y2: 0.0}

{x1: 0.0,  
y1: 1.0}  
{x2: -0.38,  
y2: 0.32}

---

Even better  
features:  
angles of  
clock hands

theta1: 45  
theta2: 0

theta1: 90  
theta2: 140

# Feature Engineering

- With good features, you will need less data and simpler models.

Raw data:  
pixel grid



---

Better  
features:  
clock hands'  
coordinates

{x1: 0.7,  
y1: 0.7}  
{x2: 0.5,  
y2: 0.0}

{x1: 0.0,  
y1: 1.0}  
{x2: -0.38,  
y2: 0.32}

---



Even better  
features:  
angles of  
clock hands

theta1: 45  
theta2: 0

theta1: 90  
theta2: 140

# Overfitting and underfitting


- **Overfitting**: network learns patterns and irregular behaviors of the training set that cannot be generalized to test/validation set. It is specific to the training set.





# Overfitting and underfitting

- **Underfitting**: you train the model more (more epochs), your model is still getting better results (lower loss) both on the training and validation sets.



# Overfitting Solution

- Underfitting is not usually a problem if the network is not too big and time consuming to train.
- How to solve the problem of overfitting?

# Overfitting Solution

- More data: not easy to find and collect.

# Overfitting Solution

- More data: not easy to find and collect.
- Smaller network.

# Overfitting Solution

- More data: not easy to find and collect.
- Smaller network.
- Add regularization.

# Overfitting Solution

- More data: not easy to find and collect.
- Smaller network.
- Add regularization.
- Add dropout.


# Smaller Network

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
model = models.Sequential()
model.add(layers.Dense(4, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(4, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

# Smaller Network





# Smaller Network

- By forcing values of a network to be smaller, you make the distribution more regular.

# Weight Regularization

- *Large weights* tend to cause *sharp transitions* in the node functions and thus *large changes* in output for *small* changes in the inputs.

# Weight Regularization


- **L1 Regularization**: The cost added is proportional to the absolute value of the weight coefficients. (L1 norm).
- **L2 Regularization**: The cost is added proportional to the square value of the weight coefficients. (AKA., weight decay).

# Weight Regularization

```
from keras import regularizers

model = models.Sequential()
model.add(layers.Dense(16, kernel_regularizer=regularizers.l2(0.001),
                        activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, kernel_regularizer=regularizers.l2(0.001),
                        activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

# Weight Regularization



# Dropout

- Remove X-percent of the data (dropout rate)
- Compromise by a scalar proportional to dropout rate.

0.3	0.2	1.5	0.0	50% dropout →	0.0	0.2	1.5	0.0	* 2
0.6	0.1	0.0	0.3		0.6	0.0	0.0	0.3	
0.2	1.9	0.3	1.2		0.0	1.9	0.3	0.0	
0.7	0.5	1.0	0.0		0.7	0.0	0.0	0.0	


# Dropout

```
model = models.Sequential()  
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))  
model.add(layers.Dropout(0.5))  
model.add(layers.Dense(16, activation='relu'))  
model.add(layers.Dropout(0.5))  
model.add(layers.Dense(1, activation='sigmoid'))
```



Dropout rate

# Dropout





# Summary

- What is your problem? What is your data? Why deep learning?
  - Image Classification

# Summary

- What is your problem? What is your data? Why deep learning?
  - Image Classification
- How do you evaluate? (what is your loss?)
  - Miss-classification

# Summary

- What is your problem? What is your data? Why deep learning?
  - Image Classification
- How do you evaluate? (what is your loss?)
  - Miss-classification
- K-fold Validation?

# Summary

- What is your problem? What is your data? Why deep learning?
  - Image Classification
- How do you evaluate? (what is your loss?)
  - Miss-classification
- K-fold Validation?
- Make/choose a simple base-line.

# Summary

- Make a model better than base-line.

# Summary

- Make a model better than base-line.
- Develop a model that overfits.
  - Make sure you can learn the training data set well and you can converge.

# Summary

- Make a model better than base-line.
- Develop a model that overfits.
  - Make sure you can learn the training data set well and you can converge.
- Use drop-out, weight regularization, change hyperparameters to get the best results.